

Machine Learning Engineer

About the Role

We are seeking a skilled Machine Learning Engineer to join our AI team. You will design, build, and deploy ML models that power our core product features, working closely with data scientists and backend engineers.

Responsibilities

- Design and implement machine learning pipelines from data ingestion to model serving.
- Train, evaluate, and fine-tune models for NLP, classification, and recommendation tasks.
- Deploy models to production using REST APIs and containerized environments.
- Monitor model performance and retrain as needed.
- Collaborate with the data team to ensure high-quality training data.

Requirements

- 3+ years of experience in machine learning or data science.
- Proficiency in Python and ML frameworks: PyTorch or TensorFlow.
- Experience with scikit-learn, Hugging Face Transformers, or similar libraries.
- Familiarity with MLflow, Weights & Biases, or other experiment tracking tools.
- Knowledge of model deployment with FastAPI, Docker, and Kubernetes.
- Strong understanding of statistics and ML fundamentals.

Nice to Have

- Experience with LLMs and prompt engineering.
- Familiarity with vector databases (Pinecone, ChromaDB).
- AWS SageMaker or GCP Vertex AI experience.

Compensation & Benefits

Competitive salary, equity package, remote-friendly, health insurance, and learning budget.